

Version Date: April 5th, 2021

Instruction Manual

NIH Database Pipeline Application (Version 2.0)

Funded by NIMH R01MH116156-01A1

**Principal Investigator: Jessica L. Nielson, PhD
Developed by: Thomas Kirsh, BS**

Table of Contents

DOWNLOADING/INSTALLING THE PIPELINE	3
GitHub REPOSITORY	3
Download Instructions.....	3
USING THE PIPELINE	3
FAKE DATASETS FOR PRACTICE	4
FEATURES OF THE APPLICATION.....	7
Main pate of the application.....	7
Collect NDA Data Dictionaries.....	8
Collect FITBIR Data Dictionaries.....	9
Preprocess NDA	9
Preprocess FITBIR	10
Plotting the Datasets (Visualize Dataset).....	11
Merge & Transform.....	15
Merge	16
Transform	17
Get Stats.....	18
Metrics	19
UPCOMING FEATURES.....	19
KNOWN BUGS.....	19
OTHER NOTES	19

Downloading/Installing the Pipeline

GitHub Repository

Download Instructions

- Location: <https://github.com/Nielson-Lab/NIH-database-pipeline>
- Click on "Releases" (underneath the "About" section) on the right side of the home page
- Find the most recent MacOS and Windows versions of the application, then click on "Assets" underneath those headers
- Only download the application ZIP file that is compatible with your system

Using the Pipeline

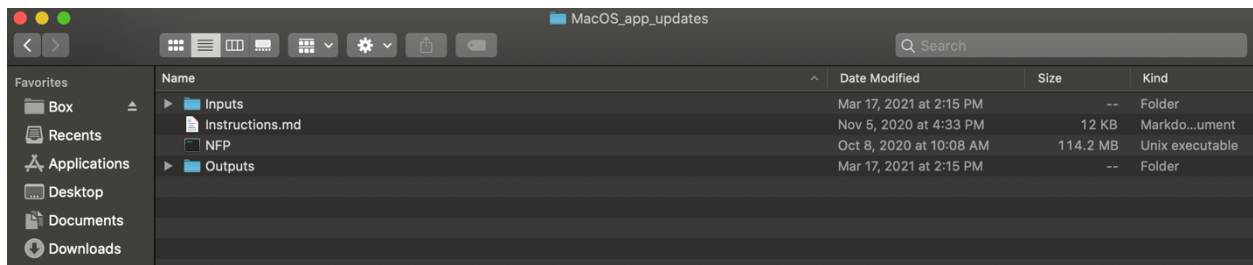


Figure 1. The main folder for the application. This folder gets downloaded from Github. Please note that although this screenshot is for the Windows version of the application, the Mac version has the same files.

The downloaded file should contain folders for *Inputs* and *Outputs*, as well as the application (NFP.exe) and an *Instructions.md* file. The user **should not** move the application out of the *MacOS_app* folder.

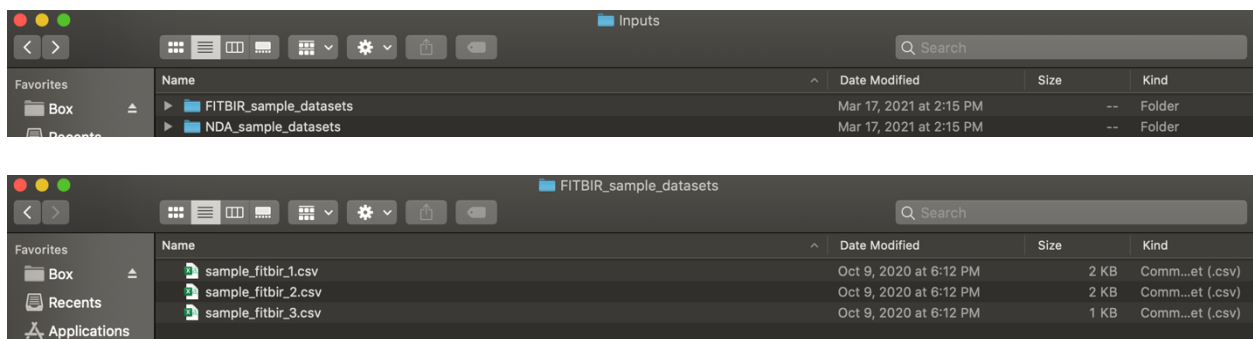


Figure 2. The *Inputs* folder in the application. Datasets that will be used in the application are uploaded here. The application only reads “.csv” or “.txt” files, so files stored in folders will not be read. Those files will need to be moved from the folder.

The application will automatically look for the files the user wants to work on in the *Inputs* folder. Therefore, the user should put the files they want to work with or on in the *Inputs* folder. The files **must** either be “.csv” or “.txt” files; “.xlsx” files are not supported at this time. Those files can easily be

FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	Info.GCSTimeOfTest
AA	69		90 Alderaan	No	0.277883	Sedation	14	Admitted				
AA	69		180 Alderaan	No	0.444051	Sedation	14	Admitted				
BB	32		90 Alderaan	No	0.435682	Other	15	Admitted				
BB	32		180 Alderaan	No	0.153481	Other	15	Admitted				
CC	33		90 Tatooine	No	0.227809	Sedation	14	Admitted				
CC	33		180 Tatooine	No	0.047977	Sedation	Untested	Admitted				
DD	65		90 Alderaan	Yes	0.637912	Paralysis	Untested	Admitted				
DD	65		180 Alderaan	Yes	0.829183	Paralysis	9	Admitted				
EE	46		90 Tatooine	No	0.940982	Other	20	Admitted				
EE	46		180 Tatooine	No	0.363848	Other	Untested	Admitted				
FF	21		90 Alderaan	Yes	0.54009	Other	Untested	Admitted				
FF	21		180 Alderaan	Yes	0.951708	Other	3	Admitted				
GG	57		90 Alderaan	No	0.122168	Sedation	5	Admitted				
GG	57		180 Alderaan	No	0.622198	Sedation	3	Admitted				
HH	58		90 Tatooine	Yes	0.383116	Paralysis	4	Admitted				
HH	58		180 Tatooine	Yes	0.628653	Paralysis	2	Admitted				
II	54		90 Alderaan	No	0.188853	Sedation	14	Admitted				
II	54		180 Alderaan	No	0.788362	Sedation	16	Admitted				
JJ	36		90 Alderaan	Yes	0.430876	Other	17	Admitted				

Figure 7. A second FITBIR sample dataset.

FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	FakeStudy	3.Info.WAISProcessSpdIndxPercntRank
AA	69	180	Blue Pill		0.23	1		
			Red Pill					
			Yellow Pill					
			Green Pill					
			Brown Pill					
			Pink Pill					
			Polka-dotted Pill					
BB	32	180	Blue Pill		0.46		2.34	
			Red Pill					
			Yellow Pill					
			Green Pill					
			Brown Pill					
			Pink Pill					
			Polka-dotted Pill					
CC	33	180	Blue Pill		0.37	0	4.56	
			Red Pill					
			Pink Pill					
			Brown Pill					
DD	65	180	Blue Pill		0.87	1		
			Red Pill					
			Pink Pill					
			Brown Pill					
EE	46	180	Blue Pill		0.64	0	9.83	

Figure 8. An example of a fake FITBIR dataset with a similar format to unflattened CSV files. The last part of the column names match an actual data element in FITBIR so the user can practice scraping data dictionaries from FITBIR.

All files that the application processes and returns will be saved to the *Outputs* folder.

Features of the Application

All files downloaded from NDA and FITBIR contain the name of the form to which they correspond. The Data Dictionary collection pages of this application rely on the downloaded dataset names to know which dictionaries to collect! Do not change the names until after you are done collecting the relevant data dictionaries.

Main Page of the Application

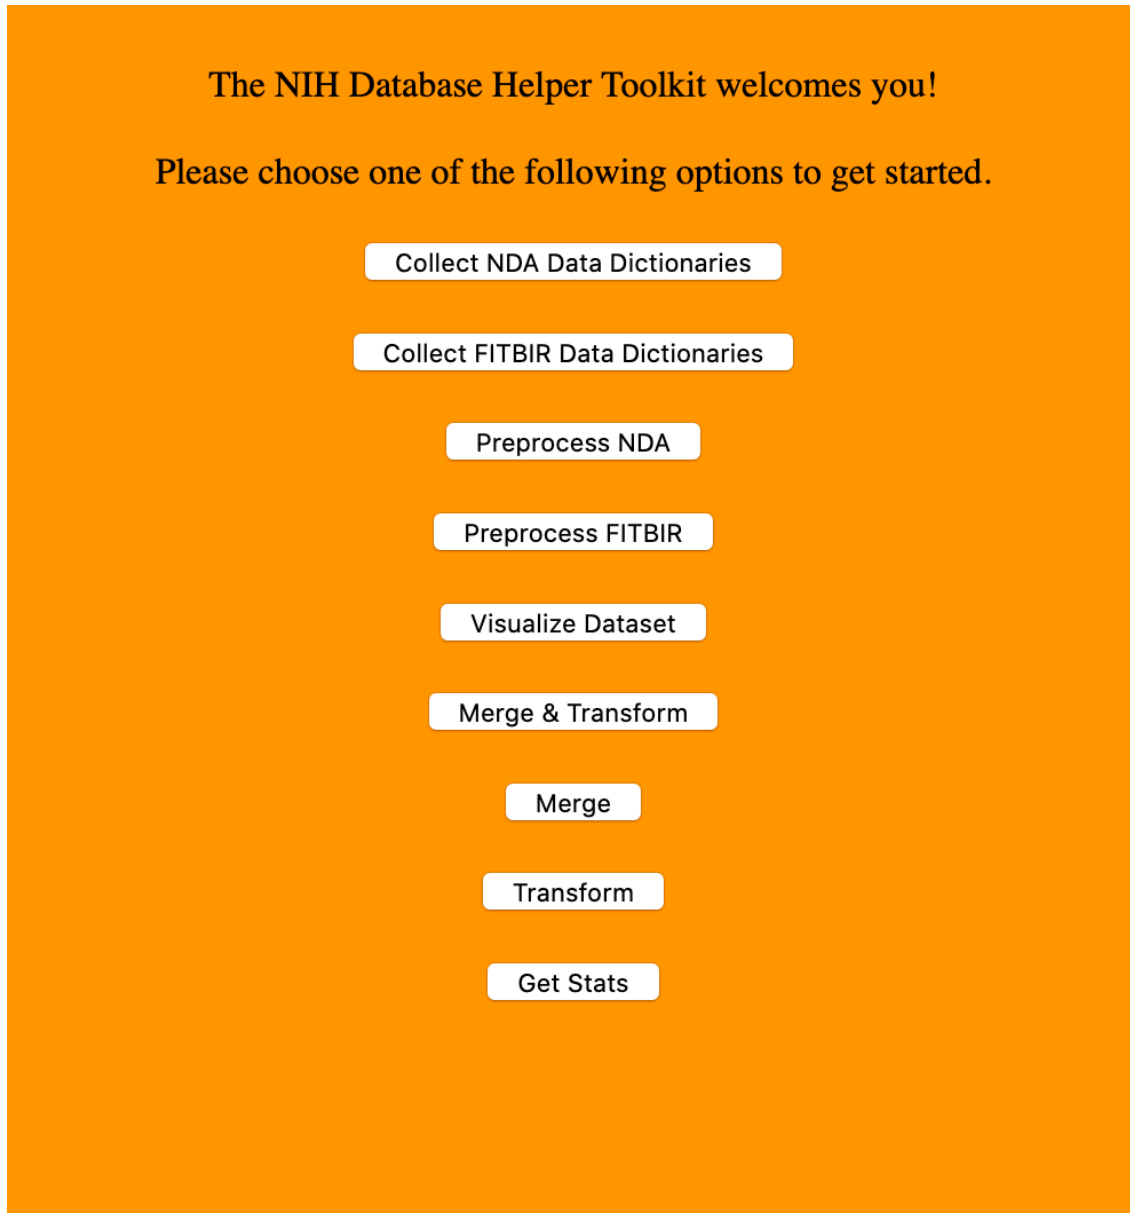


Figure 9. The main page for the web application. Each button takes the user to a different function in the application. The buttons are displayed in the order we recommend users to use the options.

The main page displays all the options the users have for working with their data. Each option describes its function. The top-down order of the options should be similar to the order the user would expect to work with their data.

Collect NDA Data Dictionaries

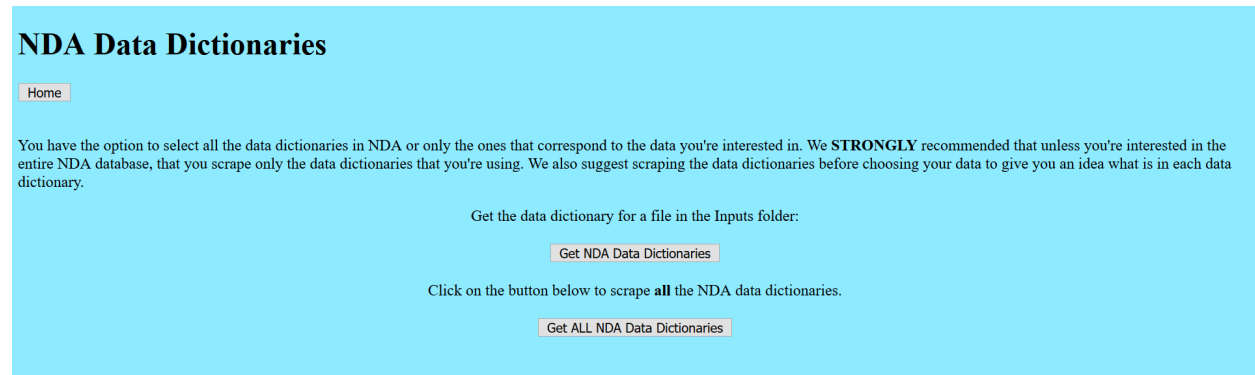
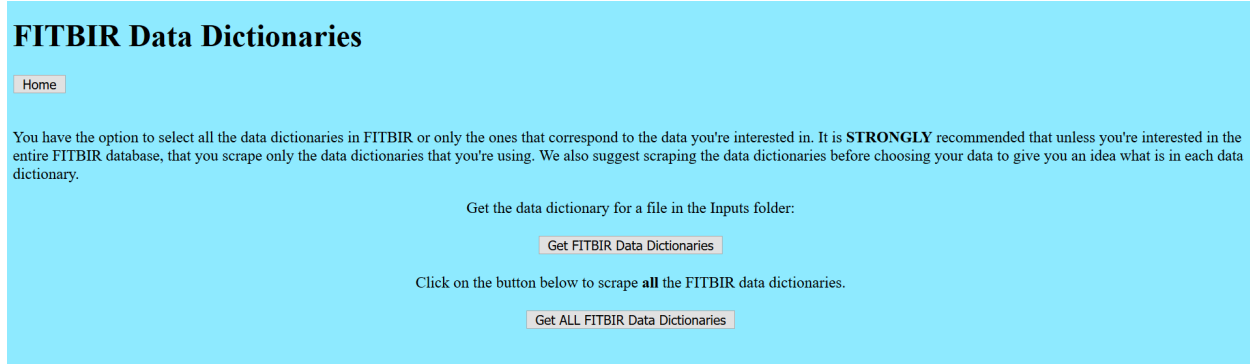


Figure 10. The page for the collecting NDA data dictionaries for all files in the *Inputs* folder. There are options to collect only the data dictionaries that correspond to the files in the *Inputs* folder and an option to collect all the data dictionaries in NDA. The second option takes more time and is useful if the user is studying the whole of NDA.

The underlying script here takes advantage of the Python API that NDA provides. Using that API, we scrape the data dictionaries, which are publicly available, that match the names of the data text files. These text files are the forms used in a study (e.g. “Beck Symptoms Inventory”). The variables are the individual *data elements* in each form (e.g. “feels sad”). From there, the script combines all the data dictionaries into one big data dictionary, with information about the variable name, its description, type, and possible values. The possible values are the values that the standard data element can have, as agreed upon by NDA, but do not necessarily reflect the actual values in your dataset. For example, the data dictionary could say that the values [1,2,3,4,5] are possible, but your dataset only has people who have values [1,2,4].

In the application, we have tried to account for when some datasets are not public. If a data dictionary cannot be found, the application will return an error. Check to make sure your files can be found in NDA’s Data Dictionary search tool.

Collect FITBIR Data Dictionaries



The screenshot shows a web page titled "FITBIR Data Dictionaries" with a light blue background. At the top left is a "Home" button. The main text explains that users can select all data dictionaries or only those corresponding to files in the *Inputs* folder, and strongly recommends scraping dictionaries before choosing data. Below this, there are two buttons: "Get FITBIR Data Dictionaries" and "Get ALL FITBIR Data Dictionaries".

Home

You have the option to select all the data dictionaries in FITBIR or only the ones that correspond to the data you're interested in. It is **STRONGLY** recommended that unless you're interested in the entire FITBIR database, that you scrape only the data dictionaries that you're using. We also suggest scraping the data dictionaries before choosing your data to give you an idea what is in each data dictionary.

Get the data dictionary for a file in the Inputs folder:

Get FITBIR Data Dictionaries

Click on the button below to scrape **all** the FITBIR data dictionaries.

Get ALL FITBIR Data Dictionaries

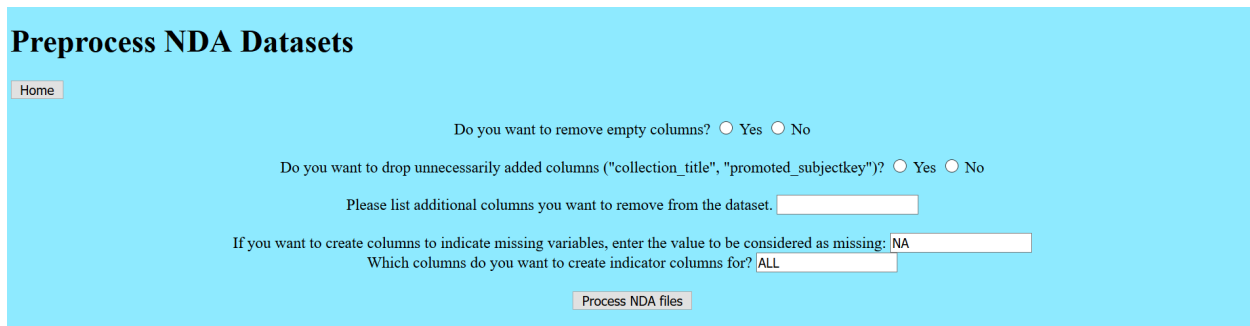
Figure 11. The page for the collecting FITBIR data dictionaries for all files in the *Inputs* folder.

There are options to collect only the data dictionaries that correspond to the files in the *Inputs* folder and an option to collect all the data dictionaries in FITBIR. The second option takes more time and is useful if the user is studying the whole of FITBIR.

In principle, the underlying script is the same as the one that collects NDA data dictionaries. The big difference is that FITBIR does not have a Python API that facilitates web scraping. Additionally, FITBIR data dictionaries (like NDA dictionaries) do not display statistics for each study in which a form was used in. This means that the value range given in FITBIR data dictionaries denotes the *possible values* that could be in this data element. A quick note about terminology: *data files* refers to the collection of forms used in a study. *forms* refer to the different assessments used in a study (i.e. Beck's Depression Inventory, PTSD Checklist, PHQ-9, etc.) while *data elements* refer to the variables measured in each assessment (i.e. "feels sad", "feels alone", etc.).

The data dictionaries can be merged with the "Get Stats" output file to create a more complete data dictionary for the user's specific dataset.

Preprocess NDA



The screenshot shows a web page titled "Preprocess NDA Datasets" with a light blue background. At the top left is a "Home" button. The page contains several form elements: two radio button questions about removing empty columns and dropping unnecessary columns, a text input for additional columns to remove, a text input for a missing value (set to "NA"), a text input for columns to create indicator columns for (set to "ALL"), and a "Process NDA files" button at the bottom.

Home

Do you want to remove empty columns? ☐ Yes ☐ No

Do you want to drop unnecessarily added columns ("collection_title", "promoted_subjectkey")? ☐ Yes ☐ No

Please list additional columns you want to remove from the dataset.

If you want to create columns to indicate missing variables, enter the value to be considered as missing:

Which columns do you want to create indicator columns for?

Process NDA files

Figure 12. The page for the user to decide how to process their NDA datasets.

Options to process NDA datafiles are:

1. Remove empty columns
2. Drop spurious columns and other columns
3. Create missing data indicator columns and for which columns.

The default missing data indicator value is an empty string (or an empty cell in Excel). If you want to tell the application that another value indicates missingness, enter that value or a list of values separated by a ‘;’. Examples of other values that indicate missingness are -777, -999, NA, NaN (as a string), etc.

Because NDA datafiles have the first row of each file as metadata, the script automatically removes them and saves those to separate files in the *Outputs* folder.

Preprocess FITBIR

Preprocess FITBIR Datasets

[Home](#)

Do you want to split all the columns? ☒ Yes ☐ No

Do you want to split some columns? ☐ Yes ☒ No

Please list the columns you want to split. Separate each column with a ‘;’

How many parts of the split column names do you wish to keep (1 means the last part will be kept, 2 means the last 2, etc.)?

Do you wish to drop certain columns? ☒ Yes ☐ No

Which columns do you want to drop? Please separate columns with a ‘;’

Do you want to remove empty columns? ☒ Yes ☐ No

If you want to create columns to indicate missing variables, enter the value to be considered as missing:

Which columns do you want to create indicator columns for? Please separate with a semi-colon:

Are you using the unflattened CSV files? ☒ Yes ☐ No

Do you want to save these columns in a separate file? ☒ Yes ☐ No

Which columns should be used to group the flatten columns? Separate each column with a ‘;’

[Process FITBIR files](#)

Figure 13. The page for the user to decide how they want to process their FITBIR files.

Options to process FITBIR files are:

1. Split the column names by period and which portion of the column name should be retained
2. Dropping specified columns. These are separated by a ‘;’.
3. Remove all empty columns.
4. Create missing value indicator columns for specific variables.
5. Fix unflattened files (by either removing the offending columns and saving them to a separate file, or by joining their values into one cell separated by a ‘;’.

Splitting the column names means separating the column names by period, so each name will be split into three parts. The option for how many parts of the name to keep tells the application to keep the last n parts. For example, if one part is requested to be kept in the column name “FakeStudy.Info.GUID”, the new column name will be “GUID”. If two parts are requested, the new column name will be “Info.GUID”. If zero or more than 3 parts are requested, the new column name will be “FakeStudy.Info.GUID”.

When downloading the datasets from FITBIR, users have the option to download the files as “flattened” or “unflattened”. We **strongly** recommend you default to downloading as “flattened”. This option creates binary variables for each value in the list columns. It is a much cleaner method for working with the data. An example of an unflattened file can be found in Figure 6.

If however, you are given an unflattened file, you can handle this in two ways:

1. Remove them and copy the GUID column and store them in a separate file.
2. Merge all the cells by group and combine the values in the list column into one string separating values by a ','.

The "group columns by" columns should be columns that don't have more than one value per row (like the list column). Good example are "GUID;AgeYrs;GeneralNotesTxt", etc.

Plotting the Datasets (Visualize Dataset)

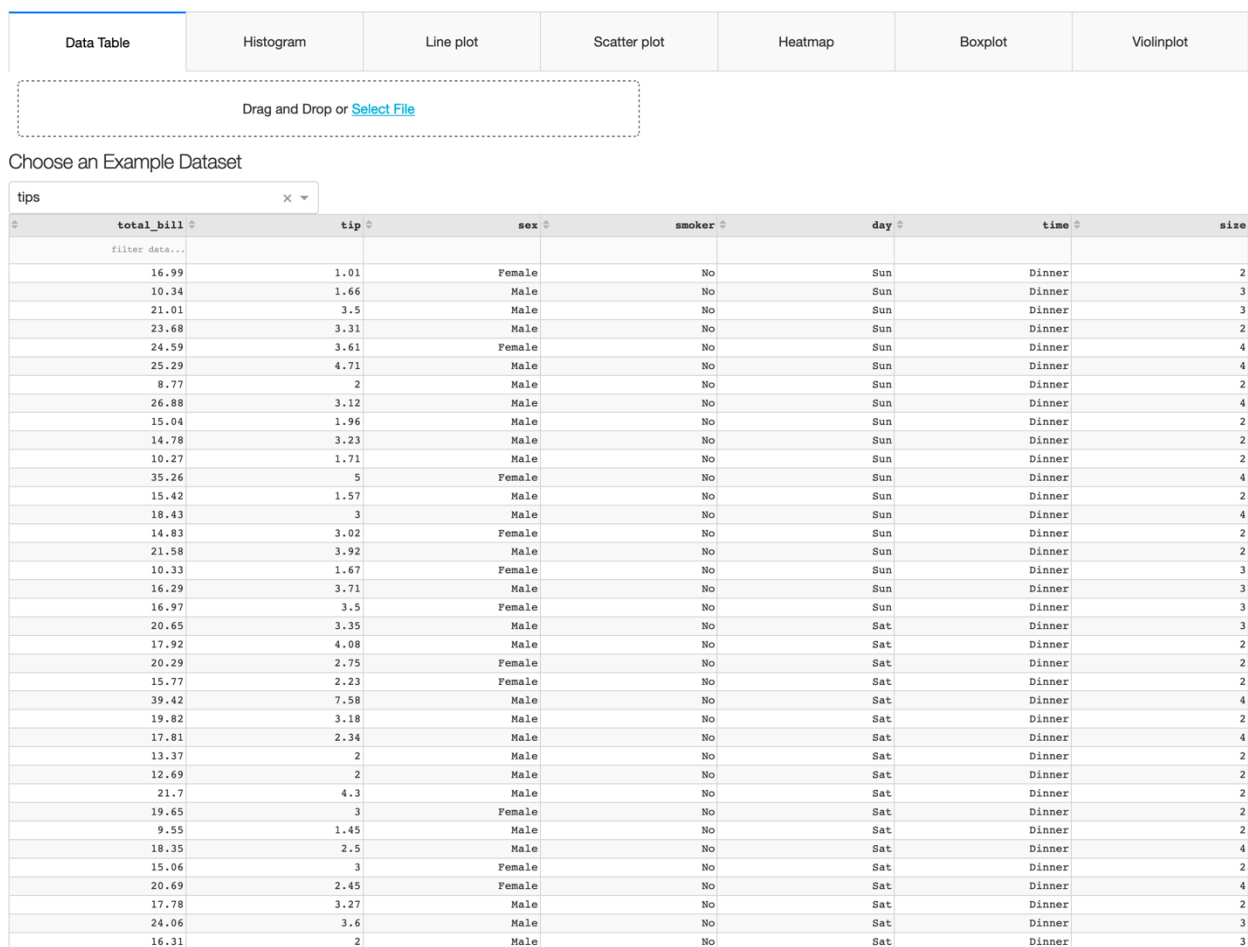


Figure 14. First page displaying the datasets in tabular form. The application provides 2 default datasets for the user to familiarize themselves with the capabilities of the visualization. The visualization page of the application no longer requires that your dataset be located in the *Inputs* folder. Instead, you can click on the “Select File” or drag and drop the file you want to look at into the dashed box. The row underneath the column names is for filtering the data. Text, numbers, and dates can be filtered. For example, you can only look at data from women by typing “Female” underneath “sex”. Using the command “= Female” (without quotes) also works. The command “contains” searches for all text values

that contain the specified substring (e.g. “contains Fem”). All columns can be filtered using “=, >, <, != (not equal), >=, <=”. Datetime columns take the format of YYYY-MM-DD (e.g. “< 2020-01” or “= 2020-01-01”). Filtering the data table does not change the dataset that is used for plots.



Figure 15. Example of a histogram for a nominal/categorical variable, plotting the count of the days people tipped, colored by “sex”. The application can produce vertical (shown here) or horizontal histograms/barplots. The counts are stacked by color (Mode can be stacked, overlay (default), or group) and not normalized (Bar Normalization can be none, fraction or percent).




Figure 16. Example of a histogram for a continuous variable. Here the “total_bill” distribution is plotted and grouped by “sex”. Instead of plotting the count for each bin, which are automatically calculated, the percent is chosen. A marginal subplot (also called a “rug plot”) is plotted above, using marks to show the

distribution of the data. The user also has the option to rename the axes labels and the title of the plot by clicking on the current text.

The histograms will only plot non-missing data. If your data has a column that has all its values missing, an error will be thrown.

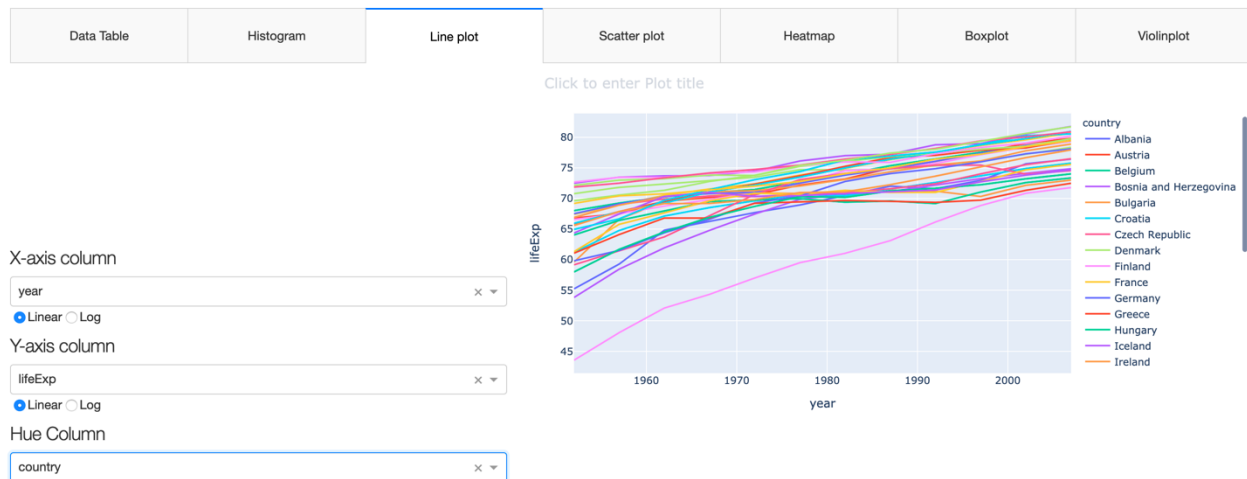


Figure 17. Example of a lineplot using the “gapminder” dataset (see the Data Table example datasets). Plots two continuous variables to search for trends.



Figure 18. Example of a scatter plot. Markers can be colored by the hue column, vary in sizes specified in the dataset, and stylized.

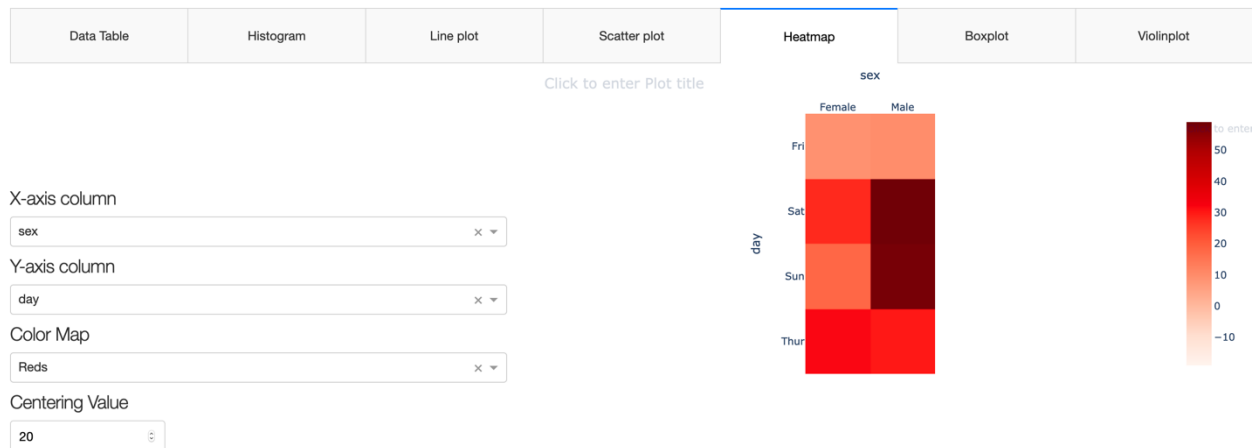


Figure 19. Example of a heatmap. Best for plotting categorical variables but numerical variables can be used as well. Multiple coloring options are included and the option to center the color scale is called “Centering Value”.

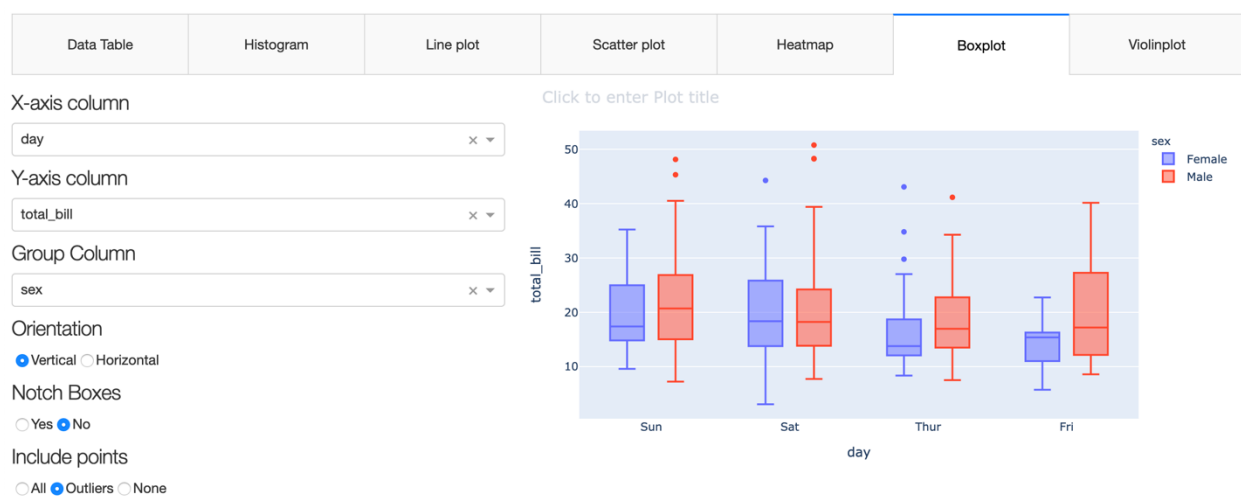


Figure 20. Example of the boxplot feature for visualizing data. These can be oriented vertically (shown) or horizontally. The columns will need to be adjusted appropriately. There is also the option to exclude outliers, show outliers, or show all data as points in addition to the boxplot.

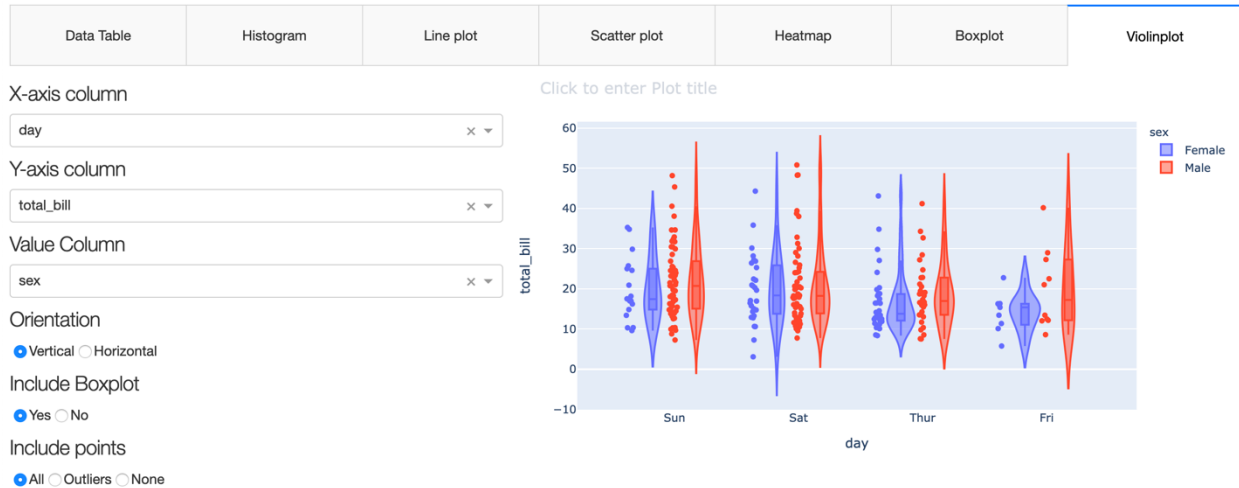


Figure 21. Example of the violin plot feature for visualizing data. These can be oriented vertically (shown) or horizontally. The columns will need to be adjusted appropriately. There is also the option to exclude outliers, show outliers, or show all data as points in addition to a boxplot.

Merge & Transform

Select which data archive your data comes from. ☐ NDA ☒ FITBIR/NIDA

Enter the name of the column that contains the GUIDs:

Enter the name of the column that contains the time points (if applicable):

Enter the columns you want to exclude from the conversion, separated by a semi-colon (;). **These should not include the GUID column or the time point column.**
For example, enter age;gender;marriage_status:

Choose the function you want to use to aggregate your data:

If the columns is a list of dates, enter the time interval in days (e.g. 30 days for a time interval of 1 month):

Do you want to use the values in the time column as your new column names? (Ex: values of '30', '90', '180' become "Col1_Day30", "Col1_Day90", etc): ☒ Yes ☐ No

Enter how you want the time points denoted (i.e. enter M if you want variables to end in M0, M6, M12, etc):

Enter the name to call the merged file:

Going to Merge and Transform

***** Collecting the names of the columns across all files *****

***** The following list of files could not be merged because they are missing one or more of the columns you tried merging on *****

***** Finished collecting the column names from all files *****

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_1.csv ...

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_2.csv ...

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_3.csv ...

***** Removing empty columns from dataset *****

***** Finished removing empty columns from dataset *****

***** Saving dataset to CSV file *****

***** Finished writing dataset to CSV file *****


Time to write CSV file is 0.00 minutes...

Figure 22. The page that prompts the user for the various inputs necessary to merge and convert files from longitudinal to wide format.

The merge and transform page merges all the files in the *Inputs* folder and then converts the resultant dataset from longitudinal to wide format. This is a streamlined process for users who know that they want to merge and transform all their data files and how they want to do it. However, each time the files are merged first and then transformed. If the user wants to transform first and then merge, they can run the 'Transform' option first from the main page and then run the 'Merge' option.

See the ‘Merge’ and ‘Transform’ options for more details.

Merge



Merge Only

[Home](#)

Select which data archive your data comes from. ☐ NDA ☒ FITBIR/NIDA

Enter the name of the column that contains the GUIDs:

Enter the name of the column that contains the time points (if applicable):

Enter the name to call the merged file:

[Merge Files](#)

Going to Merge

***** Collecting the names of the columns across all files *****

***** The following list of files could not be merged because they are missing one or more of the columns you tried merging on *****

***** Finished collecting the column names from all files *****

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_1.csv ...

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_2.csv ...

Reading in file C:\Users\kirsh012\Desktop\NFPipeline\Inputs\processed_sample_fitbir_3.csv ...

***** Removing empty columns from dataset *****

***** Finished removing empty columns from dataset *****

***** Saving dataset to CSV file *****

***** Finished writing dataset to CSV file *****

Time to write CSV file is 0.00 minutes...

***** Collecting stats for columns in dataset *****

***** Saved stats to dictionary *****

Merge complete. Open file in the 'Outputs' folder

Figure 23. The page prompts the user for the inputs necessary to merge files.

The merge option runs the same merging script that “merge and transform” uses, the only difference is that this script will output the file merged file. The files are automatically read into the script from the *Inputs* folder; the user does not need to select the files in the application.

The ‘time’ column does not have to be a ‘Time’ column (likewise the ‘GUID’ column doesn’t need to be a ‘GUID’ columns), it can be all column that you want to use to direct the merge.

If your time column has empty strings either: remove the rows with empty strings, find another column to use as your time column, or figure out what dates those empty strings indicate, or just merge on the GUID column and leave the time column option to the default value. The default value tells the application to just merge on the first column.

Transform

Convert from long to wide format

Home

Select which data archive your data comes from. ☐ NDA ☐ FITBIR/NIDA

Enter the name of the ID column in your dataset:

Enter the name of the column that contains the time points you want to flip on:

Enter the columns you want to exclude from the conversion, separated by a semi-colon (;). **These should not include the GUID column or the time point column.**
For example, enter age;gender;marriage_status:

Enter the interval between time points.
If the columns is a list of dates, enter the time interval in days (e.g. 30 days for a time interval of 1 month):

Do you want to use the values in the time column as your new column names? (Ex: values of '30', '90', '180' become "Col1_Day30", "Col1_Day90", etc): ☐ Yes ☐ No

Choose the function you want to use to aggregate your data:

mean
median
mode
first

Enter how you want the time points denoted (i.e. enter M if you want variables to end in M0, M6, M12, etc):

Enter the name you want to save the resulting file as:

Transform File

Figure 24. The page that prompts the user for the necessary inputs to convert a longitudinal dataset to a wide dataset (only for NDA and FITBIR. NIDA is another beast).

The conversion from longitudinal to wide format requires that a column that contains the time points will be used. In NDA and FITBIR (and NIDA), there are three types of ‘time’ columns: dates (like 4/3/2017), strings (‘3 months’, ‘6 months’), or numbers (30, 90, 180). The application can handle all of these. For the dates, the application converts to days from the earliest time point by default. If you have a time column with specific dates (e.g. ‘03/14/15’), sort this column from earliest to latest.

If you want the output to be in months or year or weeks, you can enter a number for the interval between time points in days (i.e. 30 days for intervals of a month), and the application will divide the raw days by the interval to convert times. To illustrate, imagine your starting date is 3/14/15 and the second time point is 6/14/15. The time interval is 3 months. If you enter the time interval as 30 (for 30 days) and a prefix of “M” (for month), the application will add “M3” to the column names. If you instead enter 90, the application will add “M1” to the column names at the second time point, because these times are the “M1” times, times before them are “M0”.

For columns with strings and/or numbers, the application can use those values as the new times. Note that this requires the user to know these things about their data prior to converting.

If your time column has empty strings either: remove the rows with empty strings, find another column to use as your time column, or figure out what dates those empty strings indicate. The application will not transform the way you expect if your column has empty strings.

Because there could be multiple measurements made at the same time point (for whatever reason), the application can aggregate over these measurements in various standard ways: using the mean, median, mode, first value, last value, or no aggregation (use only if you have one measurement per time point). The aggregation will be used to return one value per GUID.

Finally, the user can enter a prefix to denote the different time points. The default prefix is “TP”, so “TP” plus the number indicating the time will be added on to the end of each column name.

GUID	AgeYrs	SiteName	SubScore7_TP180	Col3_TP180	Col1_TP180	SubScore4_TP180	Col4_TP180	Var2_TP180	Var3_TP180	Col2_TP180	SubScore10_TP180
AA	69	Alderaan	3	Sedation	No	5	14	1	nan	0.444050615	2
BB	32	Alderaan	4	Other	No	4	15	nan	2.34	0.153480652	2
CC	33	Tatooine	5	Sedation	No	5	Untested	0	4.56	0.047976587	5
DD	65	Alderaan	3	Paralysis	Yes	4	9	1	nan	0.829183182	4
EE	46	Tatooine	1	Other	No	5	Untested	0	9.83	0.363847946	1
FF	21	Alderaan	2	Other	Yes	4	3	1	5.74	0.951708259	5
GG	57	Alderaan	3	Sedation	No	1	3	0	6.35	0.622198087	3
HH	58	Tatooine	4	Paralysis	Yes	5	2	1	3.37	0.628653195	3
II	54	Alderaan	3	Sedation	No	4	16			0.788361503	5
JJ	36	Alderaan	2	Other	No	4	13			0.356793371	3
KK	58	Naboo	2	Other	No	4	Untested			0.891385275	5
LL	46	Alderaan	5	Sedation	No	1	15			0.209004442	1
MM	64	Naboo	4	Paralysis	No	5	10			0.853744843	1
NN	55	Alderaan	1	Paralysis	No	5	10			0.637717006	5
OO	62	Naboo	4	Other	No	4	12			0.613918442	1
PP	32	Naboo	3	Sedation	No	1	13			0.614692043	2

Figure 25. Example output after using the transform options in the previous figure and the merged dataset from the ‘merge’ example. The example prefix is the default “TP”. The columns “AgeYrs” and “SiteName” were excluded from the transformation because those variables do not change over time.

Get Stats

Collect Descriptive Statistics

[Home](#)

Put file you want to collect statistics for into the Inputs folder.

Choose the file name to save the statistics to:

Get Stats

Figure 26. The page that prompts the user to decide what they want to call the file that collects various statistics for a single dataset in the *Inputs* folder.

The statistics that are calculated per column are:

- % missing
- # unique values
- Mean
- Median
- Min
- Max
- Mode
- Variance
- Standard Deviation
- 5th percentile
- 95th percentile
- Skewness
- Kurtosis

- Value Range (for up to 10 unique values)

Metrics

Source	Method	No. of Files	File Sizes	Time
Sample FITBIR datasets	Merge	3	3 KB, 1 KB, 1 KB	< 1 second
Sample FITBIR datasets	Transform	1	6 KB	< 1 second
Sample NDA datasets	Merge	3	2 KB, 5 KB, 8 KB	< 1 second
Sample NDA datasets	Transform	1	15 KB	<1 second
ABCD datasets	Merge	3	90 MB, 40 MB, 40 MB	51 minutes
Suicidality datasets	Merge	11	113 KB - 34 MB	19 minutes
Suicidality datasets	Transform	1	267 MB	13 minutes (output file size 63 KB)

Upcoming Features

- Concatenate files in the “Merge” section
- The bar, box, and violin plots will sorted by the count, alphabetically, and (for the box and violin plots, mean and median) of the categories is ascending and descending order
- Transform from wide format to longitudinal format

Known Bugs

- Data Table can't read comma-separated .txt files
- File names displayed in the “Merge”, “Transform”, and “Merge & Transform” sections are too long
- The MacOS version needs to be on Catalina 10.15.7 to use the application

Other Notes:

This application is not intended for use with imaging data.

You can merge using 1-2 columns as identifiers at a time. While the purpose of this application is to facilitate data mining and processing data from NDA, FITBIR, and NIDA, you can use datasets from other sources and use this application on them in the same way. For other datasets (or even NDA, FITBIR, or NIDA), your “GUID” column doesn’t *have* to be an actual GUID column, it can be whatever column you want to merge on. Same with the “time” column for merging. For transforming, your “time” column does need to actually be a “time” column and match formatting with NDA or FITBIR.

If you use a dataset that is not from NDA, FITBIR, or NIDA, that is comma-separated (CSV), enter the information into the application as though it were a FITBIR/NIDA dataset.